

ICP_7 : https://github.com/jxc84430/NeuralNetworks/tree/main/ICP_7

▼ Spring 2023 : CS5720

Neural Networks & Deep Learning ICP_7 : Jahnavi Chadalavada (700728443)

Deep Learning Image Classification with CNN

1. Training the model
2. Evaluating the model

Programming elements:

1. About CNN
2. Hyperparameters of CNN
3. Image classification with CNN

▼ 1. Follow the instruction below and then report how the performance changed(apply all at once)

- Convolutional input layer, 32 feature maps with a size of 3x3 and a rectifier activation function. • Dropout layer at 20%.
- Convolutional layer, 32 feature maps with a size of 3x3 and a rectifier activation function. • Max Pool layer with size 2x2.
- Convolutional layer, 64 feature maps with a size of 3x3 and a rectifier activation function. • Dropout layer at 20%.
- Convolutional layer, 64 feature maps with a size of 3x3 and a rectifier activation function. • Max Pool layer with size 2x2.
- Convolutional layer, 128 feature maps with a size of 3x3 and a rectifier activation function. • Dropout layer at 20%.
- Convolutional layer, 128 feature maps with a size of 3x3 and a rectifier activation function. • Max Pool layer with size 2x2.
- Flatten layer. • Dropout layer at 20%.
- Fully connected layer with 1024 units and a rectifier activation function. • Dropout layer at 20%.
- Fully connected layer with 512 units and a rectifier activation function. • Dropout layer at 20%.
- Fully connected output layer with 10 units and a Softmax activation function

[]

▼ Did the performance change?

Model Accuracy has reduced from 90% to 87% where as validation accuracy improved from 69% to 79%.

Execution time Increased.

Model predicted first four test images correctly with these changes where as previous code predicted one image incorrectly.

```
[10] # Simple CNN model for CIFAR-10
    import numpy
    from keras.datasets import cifar10
    from keras.models import Sequential
    from keras.layers import Dense
    from keras.layers import Dropout
    from keras.layers import Flatten
    from keras.constraints import maxnorm
    from keras.optimizers import SGD

[10] from keras.layers.convolutional import Conv2D
    from keras.layers.convolutional import MaxPooling2D
    from keras.utils import np_utils
    import matplotlib.pyplot as plt

    # fix random seed for reproducibility
    seed = 7
    numpy.random.seed(seed)

    # load data
    (X_train, y_train), (X_test, y_test) = cifar10.load_data()

    # normalize inputs from 0-255 to 0.0-1.0
    X_train = X_train.astype('float32')
    X_test = X_test.astype('float32')
    X_train = X_train / 255.0
    X_test = X_test / 255.0

    # one hot encode outputs
    y_train = np_utils.to_categorical(y_train)
    y_test = np_utils.to_categorical(y_test)
    num_classes = y_test.shape[1]

    # # Create the model
    model = Sequential()
    model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=maxnorm(3)))
    model.add(Dropout(0.2))
    model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Conv2D(64, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=maxnorm(3)))
```

```
[10] model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(128, (3, 3), input_shape=(32, 32,3), padding='same', activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=maxnorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dropout(0.2))
model.add(Dense(1024, activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

# Compile model
epochs = 25
lrate = 0.01
decay = lrate/epochs
sgd = SGD(lr=lrate, momentum=0.9, decay=decay, nesterov=False)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())

# Fit the model
history=model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)

# Final evaluation of the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))
```

```
[10] Model: "sequential_2"
-----  

Layer (type)          Output Shape       Param #
conv2d_12 (Conv2D)     (None, 32, 32, 32)    896
dropout_12 (Dropout)   (None, 32, 32, 32)    0
conv2d_13 (Conv2D)     (None, 32, 32, 32)    9248
max_pooling2d_6 (MaxPooling 2D) (None, 16, 16, 32) 0
conv2d_14 (Conv2D)     (None, 16, 16, 64)    18496
dropout_13 (Dropout)   (None, 16, 16, 64)    0
conv2d_15 (Conv2D)     (None, 16, 16, 64)    36928
max_pooling2d_7 (MaxPooling 2D) (None, 8, 8, 64) 0
conv2d_16 (Conv2D)     (None, 8, 8, 128)    73856
dropout_14 (Dropout)   (None, 8, 8, 128)    0
conv2d_17 (Conv2D)     (None, 8, 8, 128)    147584
max_pooling2d_8 (MaxPooling 2D) (None, 4, 4, 128) 0
flatten_2 (Flatten)    (None, 2048)        0
[10] dropout_15 (Dropout) (None, 2048)        0
dense_6 (Dense)        (None, 1024)        2098176
dropout_16 (Dropout)   (None, 1024)        0
dense_7 (Dense)        (None, 512)         524800
dropout_17 (Dropout)   (None, 512)         0
dense_8 (Dense)        (None, 10)          5130
-----
Total params: 2,915,114
Trainable params: 2,915,114
Non-trainable params: 0
-----  

None
Epoch 1/25
1563/1563 [=====] - 500s 318ms/step - loss: 1.8548 - accuracy: 0.3132 - val_loss: 1.5085 - val_accuracy: 0.4557
Epoch 2/25
1563/1563 [=====] - 501s 320ms/step - loss: 1.4169 - accuracy: 0.4830 - val_loss: 1.2830 - val_accuracy: 0.5396
Epoch 3/25
1563/1563 [=====] - 489s 313ms/step - loss: 1.2100 - accuracy: 0.5645 - val_loss: 1.1022 - val_accuracy: 0.6039
Epoch 4/25
1563/1563 [=====] - 492s 315ms/step - loss: 1.0610 - accuracy: 0.6237 - val_loss: 0.9752 - val_accuracy: 0.6509
Epoch 5/25
1563/1563 [=====] - 485s 310ms/step - loss: 0.9478 - accuracy: 0.6630 - val_loss: 0.8548 - val_accuracy: 0.7001
Epoch 6/25
1563/1563 [=====] - 484s 310ms/step - loss: 0.8567 - accuracy: 0.6956 - val_loss: 0.8513 - val_accuracy: 0.7006
Epoch 7/25
1563/1563 [=====] - 486s 311ms/step - loss: 0.7906 - accuracy: 0.7197 - val_loss: 0.8105 - val_accuracy: 0.7205
```

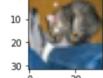
```
[10] Epoch 8/25
1563/1563 [=====] - 471s 301ms/step - loss: 0.7329 - accuracy: 0.7395 - val_loss: 0.7356 - val_accuracy: 0.7418
Epoch 9/25
1563/1563 [=====] - 466s 298ms/step - loss: 0.6885 - accuracy: 0.7586 - val_loss: 0.7237 - val_accuracy: 0.7477
Epoch 10/25
1563/1563 [=====] - 477s 305ms/step - loss: 0.6466 - accuracy: 0.7714 - val_loss: 0.7029 - val_accuracy: 0.7531
Epoch 11/25
1563/1563 [=====] - 483s 309ms/step - loss: 0.6132 - accuracy: 0.7843 - val_loss: 0.6732 - val_accuracy: 0.7604
Epoch 12/25
1563/1563 [=====] - 469s 300ms/step - loss: 0.5802 - accuracy: 0.7954 - val_loss: 0.6541 - val_accuracy: 0.7717
Epoch 13/25
1563/1563 [=====] - 482s 308ms/step - loss: 0.5554 - accuracy: 0.8035 - val_loss: 0.6496 - val_accuracy: 0.7752
Epoch 14/25
1563/1563 [=====] - 477s 305ms/step - loss: 0.5341 - accuracy: 0.8113 - val_loss: 0.6422 - val_accuracy: 0.7785
Epoch 15/25
1563/1563 [=====] - 476s 305ms/step - loss: 0.5106 - accuracy: 0.8204 - val_loss: 0.6271 - val_accuracy: 0.7851
Epoch 16/25
1563/1563 [=====] - 478s 306ms/step - loss: 0.4825 - accuracy: 0.8288 - val_loss: 0.6279 - val_accuracy: 0.7858
Epoch 17/25
1563/1563 [=====] - 476s 305ms/step - loss: 0.4630 - accuracy: 0.8345 - val_loss: 0.6322 - val_accuracy: 0.7812
Epoch 18/25
1563/1563 [=====] - 464s 297ms/step - loss: 0.4479 - accuracy: 0.8416 - val_loss: 0.6233 - val_accuracy: 0.7857
Epoch 19/25
1563/1563 [=====] - 462s 296ms/step - loss: 0.4289 - accuracy: 0.8456 - val_loss: 0.6181 - val_accuracy: 0.7906
Epoch 20/25
1563/1563 [=====] - 462s 296ms/step - loss: 0.4168 - accuracy: 0.8524 - val_loss: 0.6097 - val_accuracy: 0.7922
Epoch 21/25
1563/1563 [=====] - 478s 306ms/step - loss: 0.4023 - accuracy: 0.8576 - val_loss: 0.6375 - val_accuracy: 0.7855
Epoch 22/25
1563/1563 [=====] - 476s 305ms/step - loss: 0.3892 - accuracy: 0.8630 - val_loss: 0.6121 - val_accuracy: 0.7949
Epoch 23/25
1563/1563 [=====] - 476s 304ms/step - loss: 0.3739 - accuracy: 0.8668 - val_loss: 0.6053 - val_accuracy: 0.7941
Epoch 24/25
1563/1563 [=====] - 481s 308ms/step - loss: 0.3630 - accuracy: 0.8706 - val_loss: 0.6133 - val_accuracy: 0.7960
Epoch 25/25
750/1563 [=====>.....] - ETA: 3:58 - loss: 0.3475 - accuracy: 0.8777
```

2. Predict the first 4 images of the test data using the above model. Then, compare with the actual label for those 4 images to check whether or not the model has predicted correctly.

```
[12] # plotting first 4 images of the test data and displaying predicted output
import numpy as np
import matplotlib.pyplot as plt

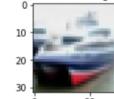
images = ["Airplane", "Automobile", "Bird", "Cat", "Deer", "Dog", "Frog", "Horse", "Ship", "Truck"]
plt.subplots_adjust(left=0.1,bottom=0.1, right=0.9,top=0.9,wspace=0.8,hspace=0.8)
for i in range(4):
    yhat=model.predict(X_test[i].reshape(1, 32, 32, 3),verbose=0)
    predicted_output=np.argmax(yhat,axis=1)
    plt.subplot(2,2,i+1)
    plt.imshow(X_test[i].reshape(32, 32,3))
    plt.title('Ground Truth of image {} : {}'.format(i+1,images[np.where(y_test[i] == 1)[0][0]]))
    plt.show()
print(f"\033[1;31m Ground truth = \033[1;34m {images[np.where(y_test[i] == 1)[0][0]]} \033[1;31m Model prediction = \033[1;35m {predicted_output}
```

Ground Truth of image 1 : Cat



Ground truth = Cat Model prediction = Cat

Ground Truth of image 2 : Ship



Ground truth = Ship Model prediction = Ship

Ground Truth of image 3 : Ship



Ground truth = Ship Model prediction = Ship

Ground Truth of image 4 : Airplane



Ground truth = Airplane Model prediction = Airplane

▼ 3. Visualize Loss and Accuracy using the history object

```
✓ [13] # plot history for accuracy
    plt.plot(history.history['accuracy'])
    plt.plot(history.history['val_accuracy'])
    plt.title('model accuracy')
    plt.ylabel('accuracy')
    plt.xlabel('epoch')
    plt.legend(['train', 'validation'], loc='lower right')
    plt.show()
# plot history for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper right')
plt.show()
```

